

Yujin Jo

Email: joyujin1015@snu.ac.kr | Phone: +82-10-9190-5510

[Google Scholar](#) | [LinkedIn](#) | [Github](#)

Location: Seoul, Republic of Korea | Expected Graduation: 02/2028

RESEARCH SUMMARY:

I am a Ph.D. candidate working on post-training of multimodal foundation models: how they acquire new capabilities without losing existing ones, and whether their reasoning stays grounded in visual evidence. At LG AI Research's EXAONE Lab, I worked on RL post-training for general-purpose VLMs across perception, grounding, and reasoning, and showed that multi-task RL does not train on the mixture one specifies: online selection turns a fixed data mixture into a hidden curriculum, and the queries that stay gradient-eligible longest are not the most useful to train on. Earlier, I developed training-free attention-level steering that keeps VLMs grounded in visual evidence rather than language priors (CVPR 2026 Findings), and memory-free continual learning that accumulates new knowledge in place without parameter growth (ICLR 2026). I am interested in bringing these questions to physical AI, where embodied reasoning models must be post-trained across heterogeneous perception and reasoning tasks while staying grounded in what they actually see.

EDUCATION:

Seoul National University, Seoul, Republic of Korea

Ph.D. Candidate in Data Science

Sep 2024 – Present

- Research Area: Reliable and efficient foundation models (Advisor: Prof. Taesup Kim)
- GPA: 4.25/4.3
- Selected Graduate Coursework (A+): Representation Learning; Training & Advancement of Foundation Models; RL, Search, and Test-Time Scaling for LLMs

M.S. in Data Science

Sep 2022 – Aug 2024

- M.S. Thesis: “Don’t Put All Knowledge in One Basket: Group-wise Prompt Learning for Vision-Language Models”
- GPA: 4.06/4.3

KAIST (Korea Advanced Institute of Science and Technology), Daejeon, Republic of Korea

B.S. in Electrical Engineering

Feb 2012 – Feb 2018

PUBLICATIONS:

- [NeurIPS 2026 Workshops] When Learning Signals Persist: Hidden Curricula in Vision–Language Reinforcement Learning | Yujin Jo, Hyunjik Jo, Yireun Kim, Jinsik Lee, Taesup Kim | [VLM4RWD](#) (Poster), [Pre-to-Post](#); extended version under review at ICLR 2027
 - Showed that zero-variance filtering in group-relative RL turns a fixed data mixture into a policy-dependent curriculum: *response demand*, a pre-RL query measure, predicts persistent gradient eligibility, and the highest-demand decile receives $2.7\text{--}3.8\times$ the accepted exposure of the lowest.
 - Showed in a controlled multi-seed intervention that queries yielding more persistent learning signal do not produce greater held-out gains, despite costing over twice the generation, and introduced a lightweight, label-free controller that corrects the resulting exposure drift.
- [ICLR 2026] Memory-Free Continual Learning with Null Space Adaptation for Zero-Shot Vision-Language Models | Yujin Jo, Taesup Kim | [arXiv:2510.21175](#)
 - Proposed Null-Space Adaptation (NuSA-CL), a memory-free continual learning method that constrains updates to under-utilized null-space directions to preserve zero-shot VLM capabilities without replay or parameter growth.
- [CVPR 2026 Findings] Attention-space Contrastive Guidance for Efficient Hallucination Mitigation in LVLMs | Yujin Jo, Sangyoon Bae, Taesup Kim | [arXiv:2601.13707](#)
 - Proposed Attention-space Contrastive Guidance (ACG), a training-free single-pass attention-level guidance method that reduces LVLM hallucinations by contrasting vision-conditioned and text-only attention paths.

- **[WACV 2025]** Retaining and Enhancing Pre-trained Knowledge in Vision-Language Models with Prompt Ensembling | Donggeun Kim*, **Yujin Jo***, Myungjoo Lee*, Taesup Kim (* equal contribution) | [arXiv:2412.07077](https://arxiv.org/abs/2412.07077)
 - Proposed an attention-isolated ensemble of learned prompt groups: train multiple prompt groups with group-wise attention masking to enforce diversity, then ensemble them at inference.

RESEARCH & WORK EXPERIENCE:

Research Intern, EXAONE Lab, LG AI Research

Mar 2026 – Aug 2026

- Worked on post-training general-purpose vision-language models, focusing on reinforcement learning for visual reasoning, perception, grounding, and chart/document understanding.
- Led the analysis of training-distribution drift in multi-task VLM RL: designed the diagnostic experiments, built the pre-RL predictor and online calibration controller, and drove the project to a first-author paper (accepted at two NeurIPS 2026 workshops; under review at ICLR 2027).

Graduate Researcher, Learning and Adaptation Algorithms Lab (LAAL), SNU

Sep 2022 – Present

- Research on VLM reliability and adaptation across the full intervention stack: decoding-time attention interventions (ACG), custom LoRA-style adapters for parameter-efficient continual adaptation (NuSA-CL), and full-parameter GRPO/RLVR post-training; experiments span CLIP, LLaVA-1.5/1.6 (LLaVA-NeXT), MiniGPT-4, Qwen-VL, Qwen2.5-VL-Instruct, and Qwen3-VL (incl. Thinking).

Teaching Assistant, Seoul National University

Mar 2023 – Dec 2025

- Designed assignments and exams, graded coursework, and provided technical support for graduate-level courses.
- Courses: Math and Statistics Foundations for Data Science (Spring 2023); Ethical Implications of AI (Spring 2025, Fall 2025).
- Received Outstanding Teaching Assistant recognition for Ethical Implications of AI (Fall 2025).

Journalist, Chosun Ilbo, Seoul, Republic of Korea

Dec 2018 – Aug 2022

- Covered technology, science, and social issues at one of South Korea's major national newspapers; interviewed domain experts and translated complex topics into accessible narratives for a broad audience.

HONORS AND AWARDS:

Lecture & Research Scholarship, Seoul National University

Spring 2023, Fall 2024

- Awarded to one graduate student per advisor per semester in Spring 2023, Fall 2024
- Full tuition waiver and stipend for research and teaching assistance

KAIST Presidential Fellowship

2012 – 2016

- Recognized for academic excellence and contributions to science and technology
- Received full scholarship for the exchange student program at INSA Toulouse (2016) and UC Berkeley (2012)

Talent Medal of Korea

Dec 2011

- Awarded by the President of the Republic of Korea to 100 outstanding young leaders annually

ACADEMIC SERVICE:

- **Reviewer:** NeurIPS 2026, ICLR 2027

SKILLS:

- **Programming & DL Frameworks:** Python, PyTorch, Hugging Face, veRL
- **Research Tools:** PEFT (LoRA, custom adapters), inference-time decoding/attention interventions, VLM evaluation, GRPO/RLVR post-training pipelines (veRL, EasyR1)
- **Languages:** Korean native; English fluent (TOEFL iBT 94, Mar 2023; IELTS 7.0, Feb 2026)